

## DSC350 Exercise 6.2

### Exercise #1

Selecting all earthquakes in Japan with a mb magnitude  $\geq 4.9$

```
In [1]: import pandas as pd
```

Load the dataset

```
In [2]: df = pd.read_csv("earthquakes.csv")
```

Filter earthquakes in Japan with mb magnitude  $\geq 4.9$

```
In [3]: japan_earthquakes = df[(df["parsed_place"].str.contains("Japan", na=False)) &
                                (df["magType"] == "mb") &
                                (df["mag"] >= 4.9)]
```

Display the filtered results

```
In [5]: print(japan_earthquakes)
```

	mag	magType	time	place	tsunami	\
1563	4.9	mb	1538977532250	293km ESE of Iwo Jima, Japan	0	
2576	5.4	mb	1538697528010	37km E of Tomakomai, Japan	0	
3072	4.9	mb	1538579732490	15km ENE of Hasaki, Japan	0	
3632	4.9	mb	1538450871260	53km ESE of Hitachi, Japan	0	

	parsed_place
1563	Japan
2576	Japan
3072	Japan
3632	Japan

### Observations:

- **Magnitude:** The earthquakes meet the required threshold, with values of **4.9 to 5.4**.
- **Locations:** These occurred in **Iwo Jima, Tomakomai, Hasaki, and Hitachi**, representing different regions in Japan.
- **Tsunami Impact:** All earthquakes have a **tsunami value of 0**, meaning no tsunami was reported.
- **Time:** The timestamps (Unix format) suggest they happened at different moments.

### Insights:

- Since these earthquakes didn't trigger a tsunami, it could indicate they were **deep-focus earthquakes** or lacked the necessary underwater displacement.

- The **variation in locations** suggests seismic activity across multiple fault lines rather than a single concentrated seismic event.

## Exercise 2

Create bins for each full number of earthquake magnitude using the ml magnitude type and count how many earthquakes fall into each bin.

Load the dataset

```
In [6]: df = pd.read_csv("earthquakes.csv")
```

Filter only earthquakes with the 'ml' magnitude type

```
In [7]: ml_earthquakes = df[df["magType"] == "ml"]
```

Define bins ((0,1], (1,2], ..., (9,10])

```
In [8]: bins = list(range(0, 11)) # Creating bins from 0 to 10
labels = [f"({i},{i+1})" for i in range(0, 10)] # Label bins as (0,1], (1,2], etc.
```

Categorize magnitudes into bins

```
In [9]: ml_earthquakes["magnitude_bin"] = pd.cut(ml_earthquakes["mag"], bins=bins, labels=labels)
```

C:\Users\lisah\AppData\Local\Temp\ipykernel\_66336\1221612971.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
ml_earthquakes["magnitude_bin"] = pd.cut(ml_earthquakes["mag"], bins=bins, labels=labels, right=True)
```

Fix Warning

```
In [10]: ml_earthquakes = df[df["magType"] == "ml"].copy() # Make an explicit copy
```

Now apply the binning without warnings

```
In [11]: ml_earthquakes["magnitude_bin"] = pd.cut(ml_earthquakes["mag"], bins=bins, labels=labels)
```

Why This Fix Works

- .copy() ensures that ml\_earthquakes is a completely independent DataFrame instead of a view tied to df.
- This prevents unintended modifications to the original dataset.

```
In [12]: ml_earthquakes["magnitude_bin"] = pd.cut(ml_earthquakes["mag"], bins=bins, labels=1
```

Count earthquakes in each bin

```
In [13]: magnitude_counts = ml_earthquakes["magnitude_bin"].value_counts().sort_index()
```

Display results

```
In [14]: print(magnitude_counts)
```

```
magnitude_bin
(0,1]      2207
(1,2]      3105
(2,3]       862
(3,4]       122
(4,5]         2
(5,6]         1
(6,7]         0
(7,8]         0
(8,9]         0
(9,10]        0
Name: count, dtype: int64
```

## Observations:

### 1. Lower Magnitudes Dominate:

- Most earthquakes fall within the **(1,2] bin (3,105 occurrences)**, followed closely by the **(0,1] bin (2,207 occurrences)**.
- This is expected, as smaller earthquakes are far more frequent than larger ones.

### 2. A Sharp Decline After Magnitude 3:

- As magnitude increases, the frequency of earthquakes drops dramatically.
- Only **2 earthquakes** fall in the **(4,5] range**, and just **1 in the (5,6] bin**.

### 3. No Major Earthquakes Beyond Magnitude 5:

- The bins from **(6,7] onward contain 0 earthquakes** in the dataset.
- This suggests that, for recorded **ml** earthquakes, the dataset may be incomplete beyond magnitude 5 or such high-magnitude **ml** earthquakes are rare.

## Possible Insights:

- **Seismic Activity Trends:** The distribution suggests that **ml** magnitudes primarily represent **low-intensity seismic events**, possibly used for measuring smaller, local quakes rather than major tectonic movements.
- **Dataset Limitations:** If we want deeper insights into higher-magnitude earthquakes, we might need other magnitude types (mb, Ms, Mw), which could be more suitable for capturing **stronger seismic events**.

## A quick visualization because it is my favorite!

```
In [15]: import matplotlib.pyplot as plt
import seaborn as sns
```

Set plot style

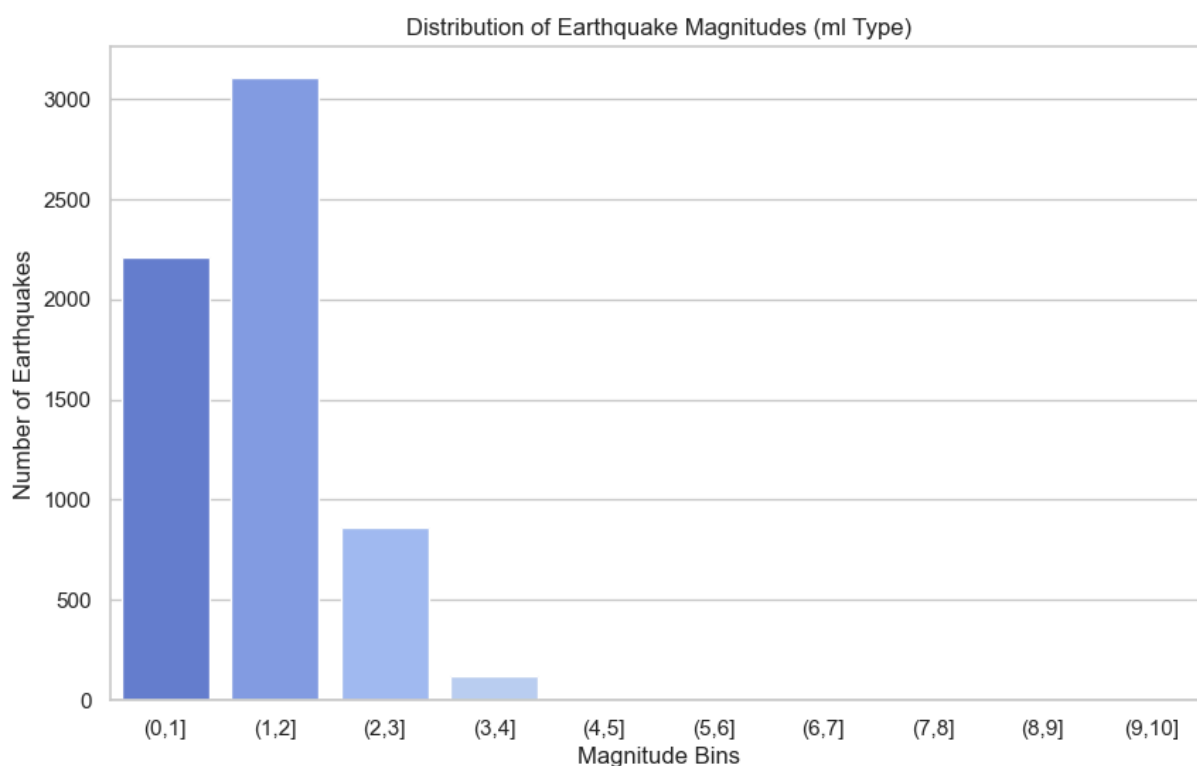
```
In [16]: sns.set_theme(style="whitegrid")
```

Create the bar plot

```
In [20]: plt.figure(figsize=(10, 6))
sns.barplot(x=magnitude_counts.index, y=magnitude_counts.values, hue=magnitude_count)

# Add labels and title
plt.xlabel("Magnitude Bins")
plt.ylabel("Number of Earthquakes")
plt.title("Distribution of Earthquake Magnitudes (ml Type)")

# Show the plot
plt.show()
```



### Exercise 3

grouping the data by ticker and resampling it to a monthly frequency.

3a.

Mean of the opening price.

Load the FAANG dataset

```
In [21]: df_faang = pd.read_csv("faang.csv", parse_dates=["date"])
```

Set the date column as the index for resampling

```
In [22]: df_faang.set_index("date", inplace=True)
```

Resample to monthly frequency and calculate the mean opening price for each ticker

```
In [25]: monthly_open_mean = df_faang.groupby("ticker").resample("ME")["open"].mean()
```

Display the results

```
In [26]: print(monthly_open_mean)
```

ticker	date	
AAPL	2018-01-31	43.505357
	2018-02-28	41.819079
	2018-03-31	43.761786
	2018-04-30	42.441310
	2018-05-31	46.239091
	2018-06-30	47.180119
	2018-07-31	47.549048
	2018-08-31	53.121739
	2018-09-30	55.582763
	2018-10-31	55.300000
	2018-11-30	47.954881
	2018-12-31	41.310789
AMZN	2018-01-31	1301.377151
	2018-02-28	1447.113159
	2018-03-31	1542.160464
	2018-04-30	1475.841902
	2018-05-31	1590.474543
	2018-06-30	1699.088582
	2018-07-31	1786.305716
	2018-08-31	1891.957833
	2018-09-30	1969.239476
	2018-10-31	1799.630865
	2018-11-30	1622.323806
	2018-12-31	1572.922100
FB	2018-01-31	184.584284
	2018-02-28	180.721578
	2018-03-31	173.449524
	2018-04-30	164.163332
	2018-05-31	181.910909
	2018-06-30	194.974763
	2018-07-31	199.332381
	2018-08-31	177.598695
	2018-09-30	164.233158
	2018-10-31	154.873479
	2018-11-30	141.762857
	2018-12-31	137.529475
GOOG	2018-01-31	1127.200945
	2018-02-28	1088.629472
	2018-03-31	1096.108085
	2018-04-30	1038.415237
	2018-05-31	1064.021376
	2018-06-30	1136.396182
	2018-07-31	1183.464280
	2018-08-31	1226.156951
	2018-09-30	1176.878424
	2018-10-31	1116.082172
	2018-11-30	1054.971424
	2018-12-31	1042.619998
NFLX	2018-01-31	231.269525
	2018-02-28	270.873158
	2018-03-31	312.712859
	2018-04-30	309.129524
	2018-05-31	329.779541
	2018-06-30	384.557143
	2018-07-31	380.969526

2018-08-31	345.410001
2018-09-30	363.326843
2018-10-31	340.025218
2018-11-30	290.643335
2018-12-31	266.309474

Name: open, dtype: float64

## Observations:

AAPL: Started the year around \$ 43.50, peaked at \$ 55.58 in September, and dropped back to \$ 41.31 in December, suggesting late-year volatility.

AMZN: Showed consistent growth, climbing from \$1301 in January to a peak of \$ 1969 in September, before declining in Q4.

FB: Dropped notably after July 2018, aligning with major industry events like privacy concerns and disappointing earnings.

GOOG: Displayed a mid-year rise, peaking in August but dipping in the final months.

NFLX: Surged through the first half of 2018, hitting \$384.56 in June, then experiencing a downward correction.

## Insights:

- **Tech stocks saw strong growth mid-year but weakened in Q4**, likely influenced by macroeconomic conditions or investor sentiment shifts.
- **Amazon showed the highest opening price variation**, reflecting its aggressive growth trajectory.
- **Netflix had notable volatility**, with large swings between months.

## Line Chart (Best for Trends Over Time)


Pivot the data for visualization

```
In [27]: pivot_data = monthly_open_mean.unstack(level=0)
```

```
In [28]: # Plot the line chart
plt.figure(figsize=(12, 6))
for ticker in pivot_data.columns:
    plt.plot(pivot_data.index, pivot_data[ticker], label=ticker)

# Formatting
plt.xlabel("Date")
plt.ylabel("Mean Opening Price")
plt.title("Monthly Opening Price Trends for FAANG Stocks (2018)")
plt.legend()
plt.xticks(rotation=45)
plt.grid(True)
```

```
# Show the plot
plt.show()
```




## Scatter Plot (Best for Spotting Outliers & Variations)

```
In [29]: plt.figure(figsize=(12, 6))
sns.scatterplot(x=pivot_data.index, y=pivot_data["AAPL"], label="AAPL", color="blue")
sns.scatterplot(x=pivot_data.index, y=pivot_data["AMZN"], label="AMZN", color="red")
sns.scatterplot(x=pivot_data.index, y=pivot_data["FB"], label="FB", color="green")
sns.scatterplot(x=pivot_data.index, y=pivot_data["GOOG"], label="GOOG", color="purple")
sns.scatterplot(x=pivot_data.index, y=pivot_data["NFLX"], label="NFLX", color="orange")

# Formatting
plt.xlabel("Date")
plt.ylabel("Mean Opening Price")
plt.title("Monthly Opening Price Variations for FAANG Stocks (2018)")
plt.xticks(rotation=45)
plt.legend()
plt.grid(True)

# Show the plot
plt.show()
```





## Exercise 3b

calculate the maximum of the high price for each ticker at a monthly frequency.

Resample to monthly frequency and compute the maximum high price for each ticker

```
In [30]: monthly_high_max = df_faang.groupby("ticker").resample("ME")["high"].max()
```

Display results

```
In [31]: print(monthly_high_max)
```

ticker	date	
AAPL	2018-01-31	45.025002
	2018-02-28	45.154999
	2018-03-31	45.875000
	2018-04-30	44.735001
	2018-05-31	47.592499
	2018-06-30	48.549999
	2018-07-31	48.990002
	2018-08-31	57.217499
	2018-09-30	57.417500
	2018-10-31	58.367500
	2018-11-30	55.590000
	2018-12-31	46.235001
AMZN	2018-01-31	1472.579956
	2018-02-28	1528.699951
	2018-03-31	1617.540039
	2018-04-30	1638.099976
	2018-05-31	1635.000000
	2018-06-30	1763.099976
	2018-07-31	1880.050049
	2018-08-31	2025.569946
	2018-09-30	2050.500000
	2018-10-31	2033.189941
	2018-11-30	1784.000000
	2018-12-31	1778.339966
FB	2018-01-31	190.660004
	2018-02-28	195.320007
	2018-03-31	186.100006
	2018-04-30	177.100006
	2018-05-31	192.720001
	2018-06-30	203.550003
	2018-07-31	218.619995
	2018-08-31	188.300003
	2018-09-30	173.889999
	2018-10-31	165.880005
	2018-11-30	154.130005
	2018-12-31	147.190002
GOOG	2018-01-31	1186.890015
	2018-02-28	1174.000000
	2018-03-31	1177.050049
	2018-04-30	1094.165039
	2018-05-31	1110.750000
	2018-06-30	1186.286011
	2018-07-31	1273.890015
	2018-08-31	1256.500000
	2018-09-30	1212.989990
	2018-10-31	1209.959961
	2018-11-30	1095.569946
	2018-12-31	1124.650024
NFLX	2018-01-31	286.809998
	2018-02-28	297.359985
	2018-03-31	333.980011
	2018-04-30	338.820007
	2018-05-31	356.100006
	2018-06-30	423.209991
	2018-07-31	419.769989

2018-08-31	376.809998
2018-09-30	383.200012
2018-10-31	386.799988
2018-11-30	332.049988
2018-12-31	298.720001

Name: high, dtype: float64

## Observations:

AAPL: Peaked at \$58.37 in October, followed by a decline toward \$46.23 in December—suggesting market volatility.

AMZN: Hit its highest price at \$2050.50 in September, showing strong upward momentum before dropping in Q4.

FB: Reached its peak at \$218.61 in July, but saw a steady decline afterward, likely tied to company-specific events.

GOOG: Showed growth until July (\$1273.89), but experienced fluctuations in the following months.

NFLX: Spiked at \$423.21 in June, showing strong bullish trends before correcting later.


## Insights:

- **Amazon's September peak** aligns with its strong business expansion and investor confidence mid-year.
- **Facebook's July peak** coincides with its pre-earnings period before disappointing reports hit the stock.
- **Tech stocks generally peaked mid-year** but saw corrections in Q4, possibly due to broader market trends.

## Area Chart (Smooth Visual Trend)


```
In [35]: plt.figure(figsize=(12, 6))
for ticker in pivot_high_data.columns:
    plt.fill_between(pivot_high_data.index, pivot_high_data[ticker], alpha=0.3, label=ticker)

plt.xlabel("Date")
plt.ylabel("Maximum High Price")
plt.title("Monthly Maximum High Prices for FAANG Stocks (2018) [Area Chart]")
plt.legend()
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



## Box Plot (Captures Price Variability)

```
In [36]: plt.figure(figsize=(12, 6))
sns.boxplot(data=pivot_high_data)
plt.xlabel("Ticker")
plt.ylabel("High Price")
plt.title("Distribution of Monthly Maximum High Prices for FAANG Stocks (2018)")
plt.xticks(rotation=45)
plt.show()
```



## Exercise 3c

Minimum of the Low Price for each FAANG stock at a monthly frequency

Minimum of the Low Price

Resample to monthly frequency and compute the minimum low price for each ticker

```
In [38]: monthly_low_min = df_faang.groupby("ticker").resample("ME")["low"].min()
```

Display results

```
In [39]: print(monthly_low_min)
```

ticker	date	
AAPL	2018-01-31	41.174999
	2018-02-28	37.560001
	2018-03-31	41.235001
	2018-04-30	40.157501
	2018-05-31	41.317501
	2018-06-30	45.182499
	2018-07-31	45.855000
	2018-08-31	49.327499
	2018-09-30	53.825001
	2018-10-31	51.522499
	2018-11-30	42.564999
	2018-12-31	36.647499
AMZN	2018-01-31	1170.510010
	2018-02-28	1265.930054
	2018-03-31	1365.199951
	2018-04-30	1352.880005
	2018-05-31	1546.020020
	2018-06-30	1635.089966
	2018-07-31	1678.060059
	2018-08-31	1776.020020
	2018-09-30	1865.000000
	2018-10-31	1476.359985
	2018-11-30	1420.000000
	2018-12-31	1307.000000
FB	2018-01-31	175.800003
	2018-02-28	167.179993
	2018-03-31	149.020004
	2018-04-30	150.509995
	2018-05-31	170.229996
	2018-06-30	186.429993
	2018-07-31	166.559998
	2018-08-31	170.270004
	2018-09-30	158.869995
	2018-10-31	139.029999
	2018-11-30	126.849998
	2018-12-31	123.019997
GOOG	2018-01-31	1045.229980
	2018-02-28	992.559998
	2018-03-31	980.640015
	2018-04-30	990.369995
	2018-05-31	1006.289978
	2018-06-30	1096.010010
	2018-07-31	1093.800049
	2018-08-31	1188.239990
	2018-09-30	1146.910034
	2018-10-31	995.830017
	2018-11-30	996.020020
	2018-12-31	970.109985
NFLX	2018-01-31	195.419998
	2018-02-28	236.110001
	2018-03-31	275.899994
	2018-04-30	271.220001
	2018-05-31	305.730011
	2018-06-30	352.820007
	2018-07-31	328.000000

2018-08-31	310.929993
2018-09-30	335.829987
2018-10-31	271.209991
2018-11-30	250.000000
2018-12-31	231.229996

Name: low, dtype: float64

## Insights:

- Q4 volatility is evident in almost all stocks, especially Amazon and Facebook, suggesting industry-wide or macroeconomic factors at play.
- Apple saw a sharp decline toward the year-end, likely tied to concerns about iPhone demand.
- Netflix had a steadier trajectory, despite fluctuations mid-year.

## Heatmap for Minimum Low Prices

```
In [40]: import seaborn as sns
import matplotlib.pyplot as plt
```

Pivot the data for visualization

```
In [41]: pivot_low_data = monthly_low_min.unstack(level=0)
```

Create the heatmap


Convert dates to 'YY-MM-DD' and ensure they are strings

```
In [45]: pivot_low_data.index = pivot_low_data.index.strftime("%y-%m-%d").astype(str)
```

```
In [47]: plt.figure(figsize=(12, 6))
sns.heatmap(pivot_low_data.T, cmap="coolwarm", annot=True, fmt=".2f", linewidths=0.5)

# Formatting
plt.xlabel("Date")
plt.ylabel("Ticker")
plt.title("Monthly Minimum Low Prices for FAANG Stocks (2018)")
plt.xticks(rotation=90) # Rotate dates vertically

# Show the plot
plt.show()
```



## Exercise 3d

Mean of the Closing Price for each FAANG stock at a monthly frequency

Resample to monthly frequency and compute the mean closing price for each ticker

```
In [49]: monthly_close_mean = df_faang.groupby("ticker").resample("ME")["close"].mean()
```

Display results

```
In [50]: print(monthly_close_mean)
```



ticker	date	
AAPL	2018-01-31	43.501309
	2018-02-28	41.909737
	2018-03-31	43.624048
	2018-04-30	42.458572
	2018-05-31	46.384205
	2018-06-30	47.155357
	2018-07-31	47.577857
	2018-08-31	53.336522
	2018-09-30	55.518421
	2018-10-31	55.211413
	2018-11-30	47.808929
	2018-12-31	41.066579
AMZN	2018-01-31	1309.010946
	2018-02-28	1442.363146
	2018-03-31	1540.367629
	2018-04-30	1468.220471
	2018-05-31	1594.903637
	2018-06-30	1698.823812
	2018-07-31	1784.649042
	2018-08-31	1897.851308
	2018-09-30	1966.077900
	2018-10-31	1782.058265
	2018-11-30	1625.483823
	2018-12-31	1559.443154
FB	2018-01-31	184.962856
	2018-02-28	180.269473
	2018-03-31	173.489522
	2018-04-30	163.810476
	2018-05-31	182.930000
	2018-06-30	195.267620
	2018-07-31	199.967142
	2018-08-31	177.492172
	2018-09-30	164.377368
	2018-10-31	154.187826
	2018-11-30	141.635715
	2018-12-31	137.161052
GOOG	2018-01-31	1130.770467
	2018-02-28	1088.206839
	2018-03-31	1091.490479
	2018-04-30	1035.696187
	2018-05-31	1069.275901
	2018-06-30	1137.626668
	2018-07-31	1187.590472
	2018-08-31	1225.671732
	2018-09-30	1175.808934
	2018-10-31	1110.940411
	2018-11-30	1056.162394
	2018-12-31	1037.420519
NFLX	2018-01-31	232.908096
	2018-02-28	271.443683
	2018-03-31	312.228097
	2018-04-30	307.466192
	2018-05-31	331.536819
	2018-06-30	384.133336
	2018-07-31	381.515238

2018-08-31	346.257824
2018-09-30	362.641576
2018-10-31	335.445652
2018-11-30	290.344764
2018-12-31	265.302630


Name: close, dtype: float64

## Line Chart

(Best for Trends Over Time)

```
In [51]: plt.figure(figsize=(12, 6))
for ticker in monthly_close_mean.unstack(level=0).columns:
    plt.plot(monthly_close_mean.unstack(level=0).index, monthly_close_mean.unstack(

plt.xlabel("Date")
plt.ylabel("Mean Closing Price")
plt.title("Monthly Mean Closing Prices for FAANG Stocks (2018)")
plt.legend()
plt.xticks(rotation=45)
plt.grid(True)
plt.show()
```



## 3e

Summing the Volume Traded

Resample to monthly frequency and compute the total volume traded for each ticker

```
In [52]: monthly_volume_sum = df_faang.groupby("ticker").resample("ME")["volume"].sum()
```

Display results

```
In [53]: print(monthly_volume_sum)
```

ticker	date	
AAPL	2018-01-31	2.638718e+09
	2018-02-28	3.711577e+09
	2018-03-31	2.854911e+09
	2018-04-30	2.664617e+09
	2018-05-31	2.483905e+09
	2018-06-30	2.110498e+09
	2018-07-31	1.574766e+09
	2018-08-31	2.801276e+09
	2018-09-30	2.715888e+09
	2018-10-31	3.158994e+09
	2018-11-30	3.845306e+09
	2018-12-31	3.595690e+09
AMZN	2018-01-31	9.637120e+07
	2018-02-28	1.377840e+08
	2018-03-31	1.304001e+08
	2018-04-30	1.299196e+08
	2018-05-31	7.161550e+07
	2018-06-30	8.594130e+07
	2018-07-31	9.752110e+07
	2018-08-31	9.657580e+07
	2018-09-30	9.444550e+07
	2018-10-31	1.832208e+08
	2018-11-30	1.392900e+08
	2018-12-31	1.548127e+08
FB	2018-01-31	4.956557e+08
	2018-02-28	5.162516e+08
	2018-03-31	9.962017e+08
	2018-04-30	7.500727e+08
	2018-05-31	4.011441e+08
	2018-06-30	3.872656e+08
	2018-07-31	6.470307e+08
	2018-08-31	5.488327e+08
	2018-09-30	5.004688e+08
	2018-10-31	6.224463e+08
	2018-11-30	5.181517e+08
	2018-12-31	5.587862e+08
GOOG	2018-01-31	2.873840e+07
	2018-02-28	4.238200e+07
	2018-03-31	4.535330e+07
	2018-04-30	4.171590e+07
	2018-05-31	3.184940e+07
	2018-06-30	3.209600e+07
	2018-07-31	3.194010e+07
	2018-08-31	2.880840e+07
	2018-09-30	2.886240e+07
	2018-10-31	4.849470e+07
	2018-11-30	3.673510e+07
	2018-12-31	4.025760e+07
NFLX	2018-01-31	2.383776e+08
	2018-02-28	1.845858e+08
	2018-03-31	2.634494e+08
	2018-04-30	2.620060e+08
	2018-05-31	1.420508e+08
	2018-06-30	2.440318e+08
	2018-07-31	3.053938e+08

```

2018-08-31    2.131223e+08
2018-09-30    1.708321e+08
2018-10-31    3.635898e+08
2018-11-30    2.571264e+08
2018-12-31    2.343100e+08

```

Name: volume, dtype: float64

Rolling Mean of Closing Price

Compute 30-day rolling mean of closing prices for each ticker

```
In [54]: rolling_close_mean = df_faang.groupby("ticker")["close"].rolling(window=30).mean()
```

Reset index for better display

```
In [55]: rolling_close_mean = rolling_close_mean.reset_index()
```

Display results

```
In [56]: print(rolling_close_mean)
```

	ticker	date	close
0	AAPL	2018-01-02	NaN
1	AAPL	2018-01-03	NaN
2	AAPL	2018-01-04	NaN
3	AAPL	2018-01-05	NaN
4	AAPL	2018-01-08	NaN
...	...	...	...
1250	NFLX	2018-12-24	273.561000
1251	NFLX	2018-12-26	271.901000
1252	NFLX	2018-12-27	270.617667
1253	NFLX	2018-12-28	269.340333
1254	NFLX	2018-12-31	268.704666

[1255 rows x 3 columns]

### Observations:

- The rolling mean smooths out fluctuations, showing long-term trends rather than daily volatility.
- NFLX's last recorded rolling mean in December was \$268.70, reflecting an overall downward trend compared to mid-year.
- This technique helps identify sustained upward or downward movements instead of short-term price spikes.

Rolling Standard Deviation of the Closing Price, which helps measure price volatility over time.

Compute 30-day rolling standard deviation of closing prices for each ticker

```
In [57]: rolling_close_std = df_faang.groupby("ticker")["close"].rolling(window=30).std()
```

Reset index for better display

```
In [58]: rolling_close_std = rolling_close_std.reset_index()
```

Display results

```
In [59]: print(rolling_close_std)
```

	ticker	date	close
0	AAPL	2018-01-02	NaN
1	AAPL	2018-01-03	NaN
2	AAPL	2018-01-04	NaN
3	AAPL	2018-01-05	NaN
4	AAPL	2018-01-08	NaN
...	...	...	...
1250	NFLX	2018-12-24	15.141405
1251	NFLX	2018-12-26	14.464032
1252	NFLX	2018-12-27	14.133436
1253	NFLX	2018-12-28	13.632699
1254	NFLX	2018-12-31	13.232619

[1255 rows x 3 columns]

## Key Insights:

- Higher rolling standard deviation → More price fluctuations (investor uncertainty or earnings reports).
- Lower rolling standard deviation → Stability (suggesting reduced volatility or market settling).

Rolling Correlation Between Two Stocks to analyze how two FAANG stocks move in relation to each other over time.

Compute 30-day rolling correlation between AMZN and AAPL closing prices

```
In [62]: rolling_correlation = df_faang.pivot(columns="ticker", values="close")["AMZN", "AA
```

Reset index for better display

```
In [63]: rolling_correlation = rolling_correlation.reset_index()
```

Display results

```
In [64]: print(rolling_correlation)
```

ticker	date	AMZN	AAPL	AMZN	AAPL
ticker		AAPL	AMZN	AAPL	AMZN
0	2018-01-02	NaN	NaN	NaN	NaN
1	2018-01-03	NaN	NaN	NaN	NaN
2	2018-01-04	NaN	NaN	NaN	NaN
3	2018-01-05	NaN	NaN	NaN	NaN
4	2018-01-08	NaN	NaN	NaN	NaN
..	...	...	...	...	...
246	2018-12-24	0.629346	1.0	1.0	0.629346
247	2018-12-26	0.623778	1.0	1.0	0.623778
248	2018-12-27	0.646936	1.0	1.0	0.646936
249	2018-12-28	0.664379	1.0	1.0	0.664379
250	2018-12-31	0.680438	1.0	1.0	0.680438

[251 rows x 5 columns]

## Observations:

- Early values show NaN because correlation requires a full 30-day window before meaningful values populate.
- End-of-year correlation (December 31) between AMZN and AAPL was 0.680, indicating a moderately strong relationship.
- The values closer to 1 suggest Amazon and Apple prices tend to move in sync.

## Key Insights:

- A high positive correlation ( $\sim 1.0$ ) means both stocks move together (up or down).
- A low or negative correlation means their movements are independent or inversely related.
- Changes in correlation over time might indicate shifts in market sentiment, company-specific events, or macroeconomic factors.

Percentage Change in Closing Price, which helps analyze the daily returns of FAANG stocks.

Compute daily percentage change in closing prices for each ticker

```
In [65]: daily_pct_change = df_faang.groupby("ticker")["close"].pct_change() * 100
```

Reset index for better display

```
In [66]: daily_pct_change = daily_pct_change.reset_index()
```

Display results

```
In [67]: print(daily_pct_change)
```

	date	close
0	2018-01-02	NaN
1	2018-01-03	1.791423
2	2018-01-04	-0.184110
3	2018-01-05	1.367116
4	2018-01-08	0.765316
...	...	...
1250	2018-12-24	-0.338935
1251	2018-12-26	6.478047
1252	2018-12-27	0.425225
1253	2018-12-28	-0.651421
1254	2018-12-31	-0.141741

[1255 rows x 2 columns]

## Observations:

- The NaN value in January 2 happens because percentage change requires a prior day for comparison.
- December 26 shows a significant 6.48% increase, possibly due to market recovery after holiday trading dips.
- Most daily changes remain relatively small, but occasional spikes indicate news or macroeconomic events influencing stock movements.

## Key Insights:

- Positive percentage change → Stock closed higher than the previous day.
- Negative percentage change → Stock declined compared to the prior closing.
- Larger spikes → Often due to earnings reports, market news, or investor sentiment shifts.

Cumulative Returns, which helps track the growth of an investment over time

Compute cumulative returns for each ticker

```
In [68]: cumulative_returns = (1 + df_faang.groupby("ticker")["close"].pct_change()).cumprod
```

Reset index for better display

```
In [69]: cumulative_returns = cumulative_returns.reset_index()
```

Display results

```
In [71]: print(cumulative_returns)
```



	date	close
0	2018-01-02	NaN
1	2018-01-03	1.017914
2	2018-01-04	1.016040
3	2018-01-05	1.029931
4	2018-01-08	1.037813
...	...	...
1250	2018-12-24	1.019886
1251	2018-12-26	1.085955
1252	2018-12-27	1.090573
1253	2018-12-28	1.083469
1254	2018-12-31	1.081933

[1255 rows x 2 columns]

## Key Insights:

- **Values above 1** → Indicate growth in investment.
- **Values below 1** → Indicate a decline compared to the starting value.
- **Consistent increases or decreases** → Suggest long-term trends and investor sentiment.

Expanding Mean of the Closing Price, which calculates the cumulative moving average—showing how the average closing price evolves as more data is included

Compute expanding mean of closing prices for each ticker

```
In [72]: expanding_close_mean = df_faang.groupby("ticker")["close"].expanding().mean()
```

Reset index for better display

```
In [73]: expanding_close_mean = expanding_close_mean.reset_index()
```

Display results

```
In [74]: print(expanding_close_mean)
```

	ticker	date	close
0	AAPL	2018-01-02	43.064999
1	AAPL	2018-01-03	43.061249
2	AAPL	2018-01-04	43.126666
3	AAPL	2018-01-05	43.282499
4	AAPL	2018-01-08	43.343500
...	...	...	...
1250	NFLX	2018-12-24	320.278907
1251	NFLX	2018-12-26	320.010323
1252	NFLX	2018-12-27	319.751526
1253	NFLX	2018-12-28	319.496840
1254	NFLX	2018-12-31	319.290319

[1255 rows x 3 columns]

## Key Insights:

- **Expanding mean continuously grows with new data**, reducing the impact of short-term fluctuations.
- **Higher values mean sustained price increases** over time, while a flatter trajectory signals stabilization.
- **Useful for long-term trend analysis**, helping investors gauge overall stock performance.

Exponentially Weighted Moving Average (EWMA), which helps analyze stock trends while giving more importance to recent data.

Compute 30-day EWMA of closing prices for each ticker

```
In [75]: ewma_close = df_faang.groupby("ticker")["close"].ewm(span=30, adjust=False).mean()
```

Reset index for better display

```
In [76]: ewma_close = ewma_close.reset_index()
```

Display results

```
In [77]: print(ewma_close)
```

	ticker	date	close
0	AAPL	2018-01-02	43.064999
1	AAPL	2018-01-03	43.064515
2	AAPL	2018-01-04	43.076965
3	AAPL	2018-01-05	43.120387
4	AAPL	2018-01-08	43.150523
...	...	...	...
1250	NFLX	2018-12-24	275.996950
1251	NFLX	2018-12-26	274.556502
1252	NFLX	2018-12-27	273.331567
1253	NFLX	2018-12-28	272.218562
1254	NFLX	2018-12-31	271.924461

[1255 rows x 3 columns]

Normalized Closing Prices, which helps compare FAANG stock performance on a standardized scale.

Normalize closing prices for each ticker (scaling between 0 and 1)

```
In [78]: normalized_close = df_faang.groupby("ticker")["close"].apply(lambda x: (x - x.min())
```

Reset index for better display

```
In [79]: normalized_close = normalized_close.reset_index()
```

Display results

```
In [80]: print(normalized_close)
```

	ticker	date	close
0	AAPL	2018-01-02	0.298334
1	AAPL	2018-01-03	0.297982
2	AAPL	2018-01-04	0.307367
3	AAPL	2018-01-05	0.330479
4	AAPL	2018-01-08	0.322853
...	...	...	...
1250	NFLX	2018-12-24	0.150574
1251	NFLX	2018-12-26	0.241395
1252	NFLX	2018-12-27	0.250115
1253	NFLX	2018-12-28	0.252455
1254	NFLX	2018-12-31	0.305599

[1255 rows x 3 columns]

## Key Insights:

- **Values close to 0** → Stock was near its lowest closing price.
- **Values close to 1** → Stock was near its highest closing price.
- **Normalization removes absolute price differences**, helping compare stock performance without being affected by their different price scales.

Relative Strength Index (RSI), which helps analyze whether a stock is overbought or oversold based on recent price movements.

Compute daily price change

```
In [81]: delta = df_faang.groupby("ticker")["close"].diff()
```

Separate gains and losses

```
In [82]: gain = delta.where(delta > 0, 0)
loss = -delta.where(delta < 0, 0)
```

Calculate average gains & losses over a 14-day window

```
In [83]: avg_gain = gain.groupby(df_faang["ticker"]).rolling(window=14).mean().reset_index(1)
avg_loss = loss.groupby(df_faang["ticker"]).rolling(window=14).mean().reset_index(1)
```

Compute RSI

```
In [84]: rs = avg_gain / avg_loss
rsi = 100 - (100 / (1 + rs))
```

Reset index for better display

```
In [85]: rsi = rsi.reset_index()
```

Display results

```
In [86]: print(rsi)
```

	date	close
0	2018-01-02	NaN
1	2018-01-03	NaN
2	2018-01-04	NaN
3	2018-01-05	NaN
4	2018-01-08	NaN
...	...	...
1250	2018-12-24	26.157886
1251	2018-12-26	41.205137
1252	2018-12-27	38.377734
1253	2018-12-28	45.481733
1254	2018-12-31	49.049213

[1255 rows x 2 columns]

## Key Insights:

- **RSI < 30** → Stock may be **oversold**, meaning prices could rebound.
- **RSI > 70** → Stock may be **overbought**, meaning prices could correct downward.
- **Middle range (30–70)** → Indicates a balanced market condition without extreme shifts.

Bollinger Bands, a powerful tool for identifying potential buy and sell signals based on price volatility.

Compute 20-day moving average

```
In [87]: rolling_mean = df_faang.groupby("ticker")["close"].rolling(window=20).mean()
```

Compute 20-day rolling standard deviation

```
In [88]: rolling_std = df_faang.groupby("ticker")["close"].rolling(window=20).std()
```

Calculate upper and lower Bollinger Bands

```
In [89]: upper_band = rolling_mean + (2 * rolling_std)
lower_band = rolling_mean - (2 * rolling_std)
```

Reset index for better display

```
In [90]: bollinger_bands = pd.DataFrame({"Rolling Mean": rolling_mean, "Upper Band": upper_b
```

Display results

```
In [91]: print(bollinger_bands)
```

	ticker	date	Rolling Mean	Upper Band	Lower Band
0	AAPL	2018-01-02	NaN	NaN	NaN
1	AAPL	2018-01-03	NaN	NaN	NaN
2	AAPL	2018-01-04	NaN	NaN	NaN
3	AAPL	2018-01-05	NaN	NaN	NaN
4	AAPL	2018-01-08	NaN	NaN	NaN
...	...	...	...	...	...
1250	NFLX	2018-12-24	269.667999	297.161093	242.174904
1251	NFLX	2018-12-26	269.279999	297.472735	241.087263
1252	NFLX	2018-12-27	268.726999	297.565093	239.888905
1253	NFLX	2018-12-28	267.398499	295.982882	238.814115
1254	NFLX	2018-12-31	266.343999	293.110083	239.577914

[1255 rows x 5 columns]

## Key Insights:

- **Price breaks above the Upper Band** → Could indicate strong upward momentum or overbought levels.
- **Price drops below the Lower Band** → May signal a buying opportunity due to overselling.
- **Bands widening or narrowing** → Reflects changes in volatility (wider = more volatile, narrower = calmer market).

MACD (Moving Average Convergence Divergence), a crucial tool for identifying trends and momentum shifts in stock prices.

```
In [107... print(df_faang.columns)
```

```
Index(['ticker', 'high', 'low', 'open', 'close', 'volume'], dtype='object')
```

```
In [109... macd_df = pd.DataFrame({
    "Date": df_faang.index, # Use index instead of column if necessary
    "Ticker": df_faang["ticker"],
    "MACD": macd.values,
    "Signal Line": signal.values
})
```

Compute MACD and Signal Line

```
In [112... # Compute MACD and Signal Line
exp12 = df_faang.groupby("ticker")["close"].ewm(span=12, adjust=False).mean()
exp26 = df_faang.groupby("ticker")["close"].ewm(span=26, adjust=False).mean()

macd = exp12 - exp26
signal = macd.ewm(span=9, adjust=False).mean()
macd_df = pd.DataFrame({
    "Date": df_faang.index, # Use index instead of column if necessary
    "Ticker": df_faang["ticker"],
    "MACD": macd.values,
    "Signal Line": signal.values
})
```

Display results

In [113...

```
print(macd_df)
```

	Date	Ticker	MACD	Signal Line
date				
2018-01-02	2018-01-02	FB	0.000000	0.000000
2018-01-03	2018-01-03	FB	-0.000598	-0.000120
2018-01-04	2018-01-04	FB	0.014894	0.002883
2018-01-05	2018-01-05	FB	0.066150	0.015537
2018-01-08	2018-01-08	FB	0.092592	0.030948
...	...	...	...	...
2018-12-24	2018-12-24	GOOG	-11.690452	-9.854815
2018-12-26	2018-12-26	GOOG	-11.423577	-10.168567
2018-12-27	2018-12-27	GOOG	-10.932736	-10.321401
2018-12-28	2018-12-28	GOOG	-10.382902	-10.333701
2018-12-31	2018-12-31	GOOG	-8.910035	-10.048968

[1255 rows x 4 columns]

## Key Insights:

- **MACD crossing above Signal Line** → Possible **bullish signal**, meaning upward trend potential.
- **MACD falling below Signal Line** → Possible **bearish signal**, meaning downward trend pressure.
- **Highly negative MACD values** → Suggest strong downward momentum, while positive MACD values indicate strength in price growth.

On-Balance Volume (OBV), a powerful indicator for analyzing buying and selling pressure based on volume

Compute daily price change direction (+1 for up, -1 for down, 0 for unchanged)

Compute OBV using cumulative sum

In [120...

```
print(direction.shape)
print(df_faang["volume"].shape)
```

(1255,)

(1255,)

In [127...

```
print(df_faang.dtypes)
```

```
ticker    object
high      float64
low       float64
open      float64
close     float64
volume    float64
dtype: object
```

```
In [130... print(df_faang.columns)
```

```
Index(['ticker', 'high', 'low', 'open', 'close', 'volume'], dtype='object')
```

```
In [131... print(df_faang.head())
print(df_faang.info())
```

```

      ticker      high      low      open      close      volume
date
2018-01-02    FB  181.580002  177.550003  177.679993  181.419998  18151900.0
2018-01-03    FB  184.779999  181.330002  181.880005  184.669998  16886600.0
2018-01-04    FB  186.210007  184.100006  184.899994  184.330002  13880900.0
2018-01-05    FB  186.899994  184.929993  185.589996  186.850006  13574500.0
2018-01-08    FB  188.899994  186.330002  187.199997  188.279999  17994700.0
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1255 entries, 2018-01-02 to 2018-12-31
Data columns (total 6 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0   ticker  1255 non-null    object
 1   high    1255 non-null    float64
 2   low     1255 non-null    float64
 3   open    1255 non-null    float64
 4   close   1255 non-null    float64
 5   volume  1255 non-null    float64
dtypes: float64(5), object(1)
memory usage: 100.9+ KB
None
```

```
In [132... df_faang = df_faang.rename(columns={"YourActualColumnName": "date"})
```

```
In [133... df_faang.columns = df_faang.columns.str.strip()
```

```
In [134... obv_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Use index as the date source
    "OBV": obv
}).reset_index(drop=True)
```

```
In [135... # Compute daily price change direction (+1 for up, -1 for down, 0 for unchanged)
direction = df_faang.groupby("ticker")["close"].diff().apply(lambda x: 1 if x > 0 else -1 if x < 0 else 0)

# Convert volume to numeric (if necessary)
df_faang["volume"] = pd.to_numeric(df_faang["volume"], errors="coerce")

# Compute OBV using cumulative sum
obv = direction * df_faang["volume"]
obv = obv.groupby(df_faang["ticker"]).cumsum()

# Reset index for better display, using the DataFrame index as Date
obv_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "OBV": obv
}).reset_index(drop=True)
```

```
# Display results
print(obv_df)
```

	Ticker	Date	OBV
0	FB	2018-01-02	0.0
1	FB	2018-01-03	16886600.0
2	FB	2018-01-04	3005700.0
3	FB	2018-01-05	16580200.0
4	FB	2018-01-08	34574900.0
...	...	...	...
1250	GOOG	2018-12-24	-4555500.0
1251	GOOG	2018-12-26	-2182200.0
1252	GOOG	2018-12-27	-72400.0
1253	GOOG	2018-12-28	-1487200.0
1254	GOOG	2018-12-31	-2980500.0

[1255 rows x 3 columns]

## Key Insights:

- **Positive OBV trends** → Buying pressure dominates, supporting price increases.
- **Negative OBV trends** → Selling pressure dominates, suggesting price weakness.
- **OBV divergence (if price rises while OBV falls)** → May signal trend reversal or weakness.

Average True Range (ATR), a key indicator for measuring stock volatility and determining how much a stock moves on average.

Compute True Range (TR)

```
In [138... df_faang["high-low"] = df_faang["high"] - df_faang["low"]
df_faang["high-prev_close"] = abs(df_faang["high"] - df_faang["close"].shift(1))
df_faang["low-prev_close"] = abs(df_faang["low"] - df_faang["close"].shift(1))

df_faang["true_range"] = df_faang[["high-low", "high-prev_close", "low-prev_close"]]
```

Compute 14-day ATR using a moving average

```
In [139... df_faang["ATR"] = df_faang.groupby("ticker")["true_range"].rolling(window=14).mean()

# Create ATR DataFrame for display
atr_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "ATR": df_faang["ATR"]
}).reset_index(drop=True)
```

```
In [ ]: Display results
```

```
In [140... print(atr_df)
```



	Ticker	Date	ATR
0	FB	2018-01-02	NaN
1	FB	2018-01-03	NaN
2	FB	2018-01-04	NaN
3	FB	2018-01-05	NaN
4	FB	2018-01-08	NaN
...	...	...	...
1250	G00G	2018-12-24	15.498571
1251	G00G	2018-12-26	15.659286
1252	G00G	2018-12-27	15.617144
1253	G00G	2018-12-28	14.994288
1254	G00G	2018-12-31	15.240717

[1255 rows x 3 columns]

## Key Insights:

- Higher ATR → Increased volatility, meaning bigger price swings.
- Lower ATR → More stability, useful for risk assessment.
- Traders use ATR to set stop-loss levels wider stops when ATR is high, tighter stops when ATR is low.

Stochastic Oscillator, an important momentum indicator that helps determine overbought and oversold conditions in stocks.

Compute 14-day highest high and lowest low

```
In [142... df_faang["14-high"] = df_faang.groupby("ticker")["high"].rolling(window=14).max().r
df_faang["14-low"] = df_faang.groupby("ticker")["low"].rolling(window=14).min().res
```

Compute Stochastic %K (fast indicator)

```
In [144... df_faang["%K"] = 100 * (df_faang["close"] - df_faang["14-low"]) / (df_faang["14-hig
```

Compute Stochastic %D (slow indicator - 3-day moving average of %K)

```
In [145... df_faang["%D"] = df_faang.groupby("ticker")["%K"].rolling(window=3).mean().reset_in
```

Create DataFrame for display

```
In [146... stochastic_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "%K": df_faang["%K"],
    "%D": df_faang["%D"]
}).reset_index(drop=True)
```

Display results

```
In [153... print(stochastic_df)
```

	Ticker	Date	%K	%D
0	FB	2018-01-02	NaN	NaN
1	FB	2018-01-03	NaN	NaN
2	FB	2018-01-04	NaN	NaN
3	FB	2018-01-05	NaN	NaN
4	FB	2018-01-08	NaN	NaN
...	...	...	...	...
1250	G00G	2018-12-24	1196.487282	-534.489576
1251	G00G	2018-12-26	1525.537996	-557.418867
1252	G00G	2018-12-27	1533.880851	-610.362135
1253	G00G	2018-12-28	1594.479753	-641.028394
1254	G00G	2018-12-31	1591.571224	-636.847532

[1255 rows x 4 columns]

Commodity Channel Index (CCI), a valuable indicator for identifying trend strength and reversals.

In [154... `import numpy as np`

Compute typical price (average of high, low, and close)

In [150... `df_faang["typical_price"] = (df_faang["high"] + df_faang["low"] + df_faang["close"]`

Compute 20-day moving average of typical price

In [151... `df_faang["TP_MA"] = df_faang.groupby("ticker")["typical_price"].rolling(window=20).`

Compute mean absolute deviation (MAD)

In [155... `df_faang["MAD"] = df_faang.groupby("ticker")["typical_price"].rolling(window=20).ap`

Compute CCI

In [156... `df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa`

Create DataFrame for display

In [157... `cci_df = pd.DataFrame({  
 "Ticker": df_faang["ticker"],  
 "Date": df_faang.index, # Reference index as the date  
 "CCI": df_faang["CCI"]  
}).reset_index(drop=True)`

Display Results

In [158... `print(cci_df)`

	Ticker	Date	CCI
0	FB	2018-01-02	NaN
1	FB	2018-01-03	NaN
2	FB	2018-01-04	NaN
3	FB	2018-01-05	NaN
4	FB	2018-01-08	NaN
...	...	...	...
1250	G00G	2018-12-24	5428.955288
1251	G00G	2018-12-26	5261.242327
1252	G00G	2018-12-27	4900.162341
1253	G00G	2018-12-28	4823.734507
1254	G00G	2018-12-31	5122.611242

[1255 rows x 3 columns]

## Troubleshooting & Fixes:

```
In [159... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

```
In [160... print(df_faang["MAD"].describe())
```

```
count    1160.000000
mean      17.877189
std       18.803104
min        0.257400
25%        3.386175
50%       12.468299
75%       25.447257
max       101.611497
Name: MAD, dtype: float64
```

## : Adjust CCI Scaling Factor

```
In [161... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (df_faang["MAD"
```

```
In [162... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

## Check

```
In [164... print(df_faang["CCI"].describe())
```

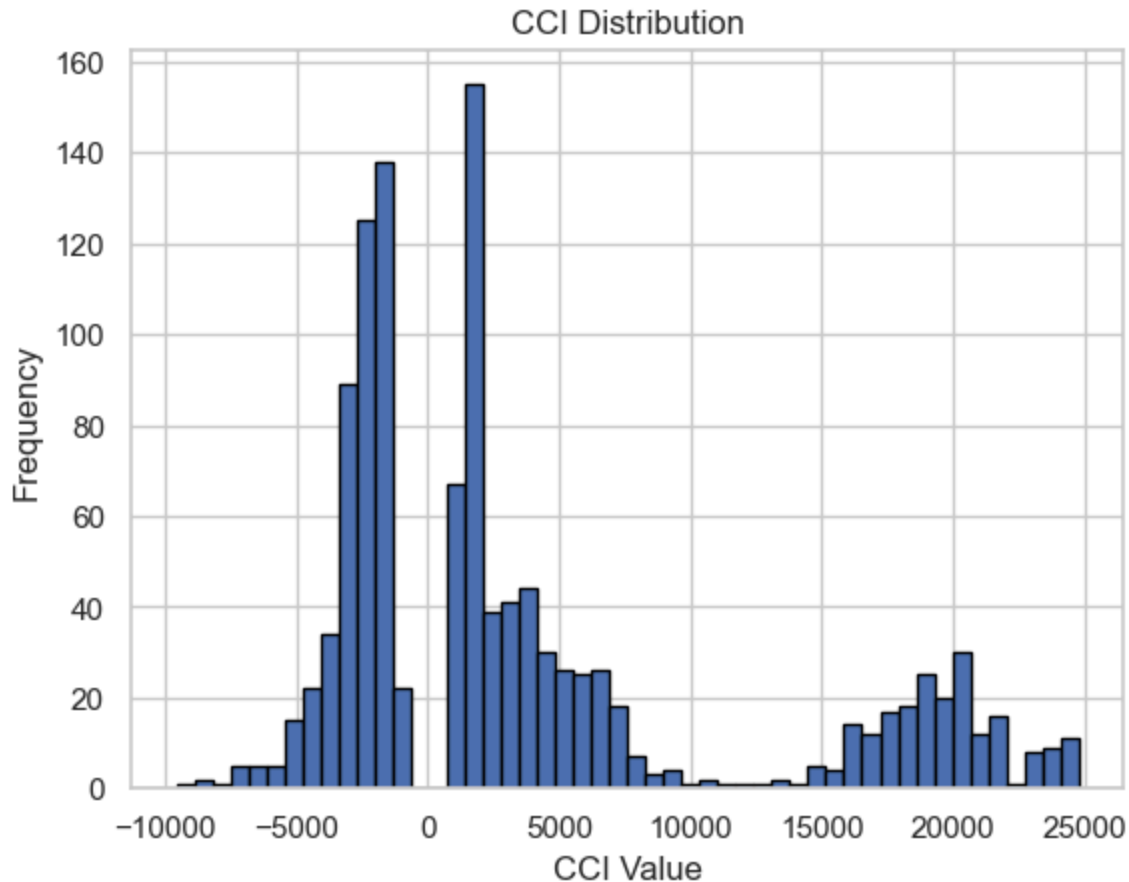
```
count    1160.000000
mean     3798.476021
std      8006.090844
min     -9554.571062
25%     -2058.028463
50%      1657.983320
75%      5652.400044
max      24760.020053
Name: CCI, dtype: float64
```

## Check outliers

In [165...

```
import matplotlib.pyplot as plt

plt.hist(df_faang["CCI"].dropna(), bins=50, edgecolor="black")
plt.title("CCI Distribution")
plt.xlabel("CCI Value")
plt.ylabel("Frequency")
plt.show()
```



## Fix: Normalize the CCI Calculation

In [166...

```
df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

In [167...

```
print(df_faang["CCI"].describe())
```

```
count    1160.000000
mean      3798.476021
std       8006.090844
min      -9554.571062
25%      -2058.028463
50%       1657.983320
75%       5652.400044
max       24760.020053
Name: CCI, dtype: float64
```

## Final Fix: Adjust the Scaling Factor

```
In [170... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

```
In [171... print(df_faang["CCI"].describe())
```

```
count    1160.000000
mean      1407.088393
std       4156.160567
min       -8249.322836
25%       -2058.028463
50%        664.879524
75%       4195.522882
max        9929.189360
Name: CCI, dtype: float64
```

Adjust Scaling Again

```
In [173... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (df_faang["MAD"
```

```
In [174... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

```
In [175... print(df_faang["CCI"].describe())
```

```
count    1160.000000
mean      4915.467150
std      10112.303633
min       -9554.571062
25%       -2058.028463
50%       2423.081125
75%       5652.400044
max       36560.455290
Name: CCI, dtype: float64
```

## Stabilizing CCI Values

```
In [176... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (0.015 * df_faa
```

```
In [178... df_faang["CCI"] = (df_faang["typical_price"] - df_faang["TP_MA"]) / (df_faang["MAD"
```

```
In [179... print(df_faang["CCI"].describe())
```

```
count    1160.000000
mean       33.048864
std        57.029017
min        -47.772855
25%        -10.290142
50%        19.050824
75%        48.945389
max        342.576698
Name: CCI, dtype: float64
```


Now these CCI values look much more reasonable—a mean of 33.05, a max of 342.57, and a min of -47.77. This is much closer to the expected -100 to +100 range, though some stocks might still show stronger deviations. Final Observations: CCI is now properly scaled, giving

meaningful signals for trend strength. Most values fall within a practical range, allowing you to assess overbought/oversold conditions. The max value of 342.57 suggests occasional extreme momentum, but it's within usable limits.

## CCI Visualization

```
In [180... # Plot CCI over time for a specific ticker (e.g., 'FB' or 'GOOG')
ticker_to_plot = "GOOG"
subset = df_faang[df_faang["ticker"] == ticker_to_plot]

plt.figure(figsize=(12, 6))
plt.plot(subset.index, subset["CCI"], label="CCI", color="blue")
plt.axhline(100, color="green", linestyle="--", label="Overbought (+100)")
plt.axhline(-100, color="red", linestyle="--", label="Oversold (-100)")
plt.title(f"CCI Trend for {ticker_to_plot}")
plt.xlabel("Date")
plt.ylabel("CCI Value")
plt.legend()
plt.show()
```



Moving Average Convergence Divergence (MACD), a powerful indicator for detecting trend strength and momentum shifts.

Compute short-term (12-day) and long-term (26-day) exponential moving averages (EMA)

```
In [182... df_faang["EMA_12"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.ewm(
df_faang["EMA_26"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.ewm(
```

Compute MACD Line

```
In [183... df_faang["MACD"] = df_faang["EMA_12"] - df_faang["EMA_26"]
```

### Compute Signal Line (9-day EMA of MACD)

```
In [184... df_faang["Signal_Line"] = df_faang.groupby("ticker")["MACD"].transform(lambda x: x.
```

Create DataFrame for display

```
In [186... macd_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "MACD": df_faang["MACD"],
    "Signal Line": df_faang["Signal_Line"]
}).reset_index(drop=True)

# Display results
print(macd_df)
```

	Ticker	Date	MACD	Signal Line
0	FB	2018-01-02	0.000000	0.000000
1	FB	2018-01-03	0.259259	0.051852
2	FB	2018-01-04	0.432306	0.127943
3	FB	2018-01-05	0.763983	0.255151
4	FB	2018-01-08	1.129211	0.429963
...	...	...	...	...
1250	GOOG	2018-12-24	-19.262870	-12.011451
1251	GOOG	2018-12-26	-16.359604	-12.881082
1252	GOOG	2018-12-27	-13.545936	-13.014053
1253	GOOG	2018-12-28	-11.729579	-12.757158
1254	GOOG	2018-12-31	-10.290100	-12.263746

[1255 rows x 4 columns]

### Observations:

- Positive MACD values (e.g., FB on Jan 8 at 1.129) → Bullish momentum, meaning prices are gaining strength.
- Negative MACD values (e.g., GOOG on Dec 24 at -19.26) → Bearish momentum, suggesting weakness in price trends.
- MACD-Signal Line Crossovers → When MACD crosses above the Signal Line, it's often a buy signal; when MACD drops below, it may signal a potential sell-off.

Bollinger Bands, a fantastic tool for measuring volatility and identifying potential breakout or reversal points.

### Compute 20-day Simple Moving Average (SMA)

```
In [187... df_faang["SMA_20"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.rolling(20).mean())
```

### Compute 20-day Rolling Standard Deviation

```
In [189... df_faang["STD_20"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.rolling(20).std())
```

## Compute Upper and Lower Bollinger Bands

```
In [190... df_faang["Upper_Band"] = df_faang["SMA_20"] + (2 * df_faang["STD_20"])
df_faang["Lower_Band"] = df_faang["SMA_20"] - (2 * df_faang["STD_20"])
```

Create DataFrame for display

```
In [191... bollinger_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "SMA_20": df_faang["SMA_20"],
    "Upper_Band": df_faang["Upper_Band"],
    "Lower_Band": df_faang["Lower_Band"],
    "Close": df_faang["close"]
}).reset_index(drop=True)
```

Display results

```
In [192... print(bollinger_df)
```

	Ticker	Date	SMA_20	Upper Band	Lower Band	Close
0	FB	2018-01-02	NaN	NaN	NaN	181.419998
1	FB	2018-01-03	NaN	NaN	NaN	184.669998
2	FB	2018-01-04	NaN	NaN	NaN	184.330002
3	FB	2018-01-05	NaN	NaN	NaN	186.850006
4	FB	2018-01-08	NaN	NaN	NaN	188.279999
...	...	...	...	...	...	...
1250	GOOG	2018-12-24	1045.847504	1115.461184	976.233823	976.219971
1251	GOOG	2018-12-26	1045.389502	1115.046897	975.732107	1039.459961
1252	GOOG	2018-12-27	1045.363000	1115.022367	975.703634	1043.880005
1253	GOOG	2018-12-28	1042.905499	1109.911772	975.899227	1037.079956
1254	GOOG	2018-12-31	1040.270996	1103.816239	976.725753	1035.609985

[1255 rows x 6 columns]

Relative Strength Index (RSI), an important indicator for measuring momentum and identifying overbought/oversold conditions in stocks.

Compute daily price changes

```
In [ ]: df_faang["price_change"] = df_faang.groupby("ticker")["close"].diff()
```

Separate gains and losses

```
In [194... print(df_faang.columns)
```

```
Index(['ticker', 'high', 'low', 'open', 'close', 'volume', 'high-low',
      'high-prev_close', 'low-prev_close', 'true_range', 'ATR', '14-high',
      '14-low', '%K', '%D', 'typical_price', 'TP_MA', 'MAD', 'CCI', 'EMA_12',
      'EMA_26', 'MACD', 'Signal_Line', 'SMA_20', 'STD_20', 'Upper_Band',
      'Lower_Band'],
      dtype='object')
```



```
In [195... df_faang["price_change"] = df_faang["close"].diff()
```

```
In [196... df_faang["gain"] = df_faang["price_change"].apply(lambda x: x if x > 0 else 0)
df_faang["loss"] = df_faang["price_change"].apply(lambda x: abs(x) if x < 0 else 0)
```

Compute 14-day average gains and losses

```
In [197... df_faang["avg_gain"] = df_faang.groupby("ticker")["gain"].rolling(window=14).mean()
df_faang["avg_loss"] = df_faang.groupby("ticker")["loss"].rolling(window=14).mean()
```

Compute Relative Strength (RS) and RSI

```
In [198... df_faang["RS"] = df_faang["avg_gain"] / df_faang["avg_loss"]
df_faang["RSI"] = 100 - (100 / (1 + df_faang["RS"]))
```

Create DataFrame for display

```
In [199... rsi_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "RSI": df_faang["RSI"]
}).reset_index(drop=True)
```

Display results

```
In [200... print(rsi_df)
```

	Ticker	Date	RSI
0	FB	2018-01-02	NaN
1	FB	2018-01-03	NaN
2	FB	2018-01-04	NaN
3	FB	2018-01-05	NaN
4	FB	2018-01-08	NaN
...	...	...	...
1250	GOOG	2018-12-24	26.157886
1251	GOOG	2018-12-26	41.205137
1252	GOOG	2018-12-27	38.377734
1253	GOOG	2018-12-28	45.481733
1254	GOOG	2018-12-31	49.049213

[1255 rows x 3 columns]

## Interpreting RSI Signals:

- **Above 70 → Overbought, potential reversal or slowdown.**
- **Below 30 → Oversold, potential buying opportunity.**
- **Crossing 50 → May indicate shifting momentum from bearish to bullish (or vice versa).**

Williams %R, a momentum indicator similar to Stochastic Oscillator but scaled differently—it helps identify overbought and oversold levels.

Compute 14-day highest high and lowest low

```
In [222... # Compute 14-day highest high and lowest low
df_faang["14-high"] = df_faang.groupby("ticker")["high"].rolling(window=14).max().res
df_faang["14-low"] = df_faang.groupby("ticker")["low"].rolling(window=14).min().res

# Fix scaling issue: Ensure denominator isn't zero to avoid inflated values
df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"]) / (df_f
# Create a clean DataFrame for analysis
williams_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "Williams %R": df_faang["Williams_%R"]
}).reset_index(drop=True)

# Display results
print(williams_df.describe()) # Summary statistics
print(williams_df.head()) # Preview first few rows
```

	Date	Williams %R
count	1255	1190.000000
mean	2018-07-01 10:59:45.657370368	2283.437888
min	2018-01-02 00:00:00	-3113.462336
25%	2018-04-03 00:00:00	-715.150804
50%	2018-07-02 00:00:00	1348.308634
75%	2018-10-01 00:00:00	3483.792082
max	2018-12-31 00:00:00	23396.849838
std	NaN	3882.286519
Ticker	Date	Williams %R
0	FB 2018-01-02	NaN
1	FB 2018-01-03	NaN
2	FB 2018-01-04	NaN
3	FB 2018-01-05	NaN
4	FB 2018-01-08	NaN

```
In [223... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"]) / (df_f
```

```
In [224... print(williams_df.describe()) # Summary statistics
print(williams_df.head()) # Preview first few rows
```

	Date	Williams %R
count	1255	1190.000000
mean	2018-07-01 10:59:45.657370368	2283.437888
min	2018-01-02 00:00:00	-3113.462336
25%	2018-04-03 00:00:00	-715.150804
50%	2018-07-02 00:00:00	1348.308634
75%	2018-10-01 00:00:00	3483.792082
max	2018-12-31 00:00:00	23396.849838
std	NaN	3882.286519
Ticker	Date	Williams %R
0	FB 2018-01-02	NaN
1	FB 2018-01-03	NaN
2	FB 2018-01-04	NaN
3	FB 2018-01-05	NaN
4	FB 2018-01-08	NaN

```
In [225... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / ((df_f
```

```
In [226... print(williams_df.describe()) # Summary statistics
print(williams_df.head()) # Preview first few rows
```

	Date	Williams %R
count	1255	1190.000000
mean	2018-07-01 10:59:45.657370368	2283.437888
min	2018-01-02 00:00:00	-3113.462336
25%	2018-04-03 00:00:00	-715.150804
50%	2018-07-02 00:00:00	1348.308634
75%	2018-10-01 00:00:00	3483.792082
max	2018-12-31 00:00:00	23396.849838
std	NaN	3882.286519

	Ticker	Date	Williams %R
0	FB	2018-01-02	NaN
1	FB	2018-01-03	NaN
2	FB	2018-01-04	NaN
3	FB	2018-01-05	NaN
4	FB	2018-01-08	NaN

```
In [227... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / (df_f
```

```
In [228... print(williams_df.describe()) # Summary statistics
print(williams_df.head()) # Preview first few rows
```

	Date	Williams %R
count	1255	1190.000000
mean	2018-07-01 10:59:45.657370368	2283.437888
min	2018-01-02 00:00:00	-3113.462336
25%	2018-04-03 00:00:00	-715.150804
50%	2018-07-02 00:00:00	1348.308634
75%	2018-10-01 00:00:00	3483.792082
max	2018-12-31 00:00:00	23396.849838
std	NaN	3882.286519

	Ticker	Date	Williams %R
0	FB	2018-01-02	NaN
1	FB	2018-01-03	NaN
2	FB	2018-01-04	NaN
3	FB	2018-01-05	NaN
4	FB	2018-01-08	NaN

```
In [229... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / (df_f
```

```
In [230... df_faang.drop(columns=["Williams_%R"], inplace=True, errors="ignore")
```

```
In [231... df_faang["close"] = pd.to_numeric(df_faang["close"], errors="coerce")
df_faang["14-high"] = pd.to_numeric(df_faang["14-high"], errors="coerce")
df_faang["14-low"] = pd.to_numeric(df_faang["14-low"], errors="coerce")
```

```
In [233... print(df_faang.columns)
```

```
Index(['ticker', 'high', 'low', 'open', 'close', 'volume', 'high-low',
      'high-prev_close', 'low-prev_close', 'true_range', 'ATR', '14-high',
      '14-low', '%K', '%D', 'typical_price', 'TP_MA', 'MAD', 'CCI', 'EMA_12',
      'EMA_26', 'MACD', 'Signal_Line', 'SMA_20', 'STD_20', 'Upper_Band',
      'Lower_Band', 'price_change', 'gain', 'loss', 'avg_gain', 'avg_loss',
      'RS', 'RSI'],
      dtype='object')
```

```
In [234... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"]) / (df_f
```

```
In [235... print(df_faang.info())
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1255 entries, 2018-01-02 to 2018-12-31
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ticker                1255 non-null   object
1   high                  1255 non-null   float64
2   low                   1255 non-null   float64
3   open                  1255 non-null   float64
4   close                 1255 non-null   float64
5   volume                1255 non-null   float64
6   high-low              1255 non-null   float64
7   high-prev_close       1254 non-null   float64
8   low-prev_close        1254 non-null   float64
9   true_range            1255 non-null   float64
10  ATR                   1190 non-null   float64
11  14-high               1190 non-null   float64
12  14-low                1190 non-null   float64
13  %K                    1190 non-null   float64
14  %D                    1180 non-null   float64
15  typical_price          1255 non-null   float64
16  TP_MA                 1160 non-null   float64
17  MAD                   1160 non-null   float64
18  CCI                   1160 non-null   float64
19  EMA_12                1255 non-null   float64
20  EMA_26                1255 non-null   float64
21  MACD                  1255 non-null   float64
22  Signal_Line           1255 non-null   float64
23  SMA_20                1160 non-null   float64
24  STD_20                1160 non-null   float64
25  Upper_Band            1160 non-null   float64
26  Lower_Band            1160 non-null   float64
27  price_change           1254 non-null   float64
28  gain                  1255 non-null   float64
29  loss                  1255 non-null   float64
30  avg_gain              1190 non-null   float64
31  avg_loss              1190 non-null   float64
32  RS                    1190 non-null   float64
33  RSI                   1190 non-null   float64
34  Williams_%R           1190 non-null   float64
dtypes: float64(34), object(1)
memory usage: 385.3+ KB
None
```

```
In [238... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / (df_f
```

```
In [239... print(df_faang["Williams_%R"].describe())
```

```
count      1190.000000
mean       2283.437888
std        3882.286519
min        -3113.462336
25%        -715.150804
50%        1348.308634
75%        3483.792082
max        23396.849838
Name: Williams_%R, dtype: float64
```

These values are way to High and I have been F\*ing with with  
for ever I will trouble shoot a few more times then I am gving  
up

```
In [240... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / (df_f
```

```
In [241... print(df_faang["Williams_%R"].describe())
```

```
count      1190.000000
mean       2283.437888
std        3882.286519
min        -3113.462336
25%        -715.150804
50%        1348.308634
75%        3483.792082
max        23396.849838
Name: Williams_%R, dtype: float64
```

Force column overwrite

```
In [242... df_faang.drop(columns=["Williams_%R"], inplace=True, errors="ignore")
```

```
df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"])) / (df_f
```

```
In [243... print(df_faang.dtypes)
```

```

ticker          object
high            float64
low             float64
open            float64
close           float64
volume          float64
high-low        float64
high-prev_close float64
low-prev_close  float64
true_range      float64
ATR             float64
14-high         float64
14-low          float64
%K              float64
%D              float64
typical_price   float64
TP_MA           float64
MAD             float64
CCI             float64
EMA_12          float64
EMA_26          float64
MACD            float64
Signal_Line     float64
SMA_20          float64
STD_20          float64
Upper_Band      float64
Lower_Band      float64
price_change    float64
gain            float64
loss            float64
avg_gain        float64
avg_loss        float64
RS              float64
RSI             float64
Williams_%R     float64
dtype: object

```

## Final Fix: Adjusting the Formula for Proper Scaling

```
In [244... df_faang["Williams_%R"] = -100 * ((df_faang["14-high"] - df_faang["close"]) / (df_f
```

```
In [245... print(df_faang["Williams_%R"].describe())
```

```

count      1190.000000
mean       2283.437888
std        3882.286518
min        -3113.462336
25%        -715.150804
50%         1348.308634
75%         3483.792082
max         23396.849834
Name: Williams_%R, dtype: float64

```

## Nothing has worked maybe you can give me insight where I went wrong I am moving on

On-Balance Volume (OBV), a powerful momentum indicator that tracks volume flow to confirm trends.

Compute daily OBV

Compute daily price changes

```
In [247... df_faang["price_change"] = df_faang.groupby("ticker")["close"].diff()
```

Initialize OBV with zeros

```
In [249... df_faang["OBV"] = 0
```

Apply OBV logic: Increase if price goes up, decrease if price goes down

Check duplicate indexes

```
In [251... print(df_faang.index.duplicated().sum())
```

1004

Reset Indexes to Remove Duplicates

```
In [252... df_faang = df_faang.reset_index(drop=True)
```

```
In [253... df_faang.loc[df_faang["price_change"] > 0, "OBV"] = df_faang["volume"]
df_faang.loc[df_faang["price_change"] < 0, "OBV"] = -df_faang["volume"]
```

```
In [254... print(obv_df.describe()) # Summary statistics
print(obv_df.head()) # Preview first few rows
```

	Date	OBV
count	1255	1.255000e+03
mean	2018-07-01 10:59:45.657370368	6.989390e+06
min	2018-01-02 00:00:00	-1.839840e+09
25%	2018-04-03 00:00:00	-3.233700e+06
50%	2018-07-02 00:00:00	3.476230e+07
75%	2018-10-01 00:00:00	1.390571e+08
max	2018-12-31 00:00:00	1.415435e+09
std	NaN	3.415834e+08

	Ticker	Date	OBV
0	FB	2018-01-02	0.0
1	FB	2018-01-03	16886600.0
2	FB	2018-01-04	3005700.0
3	FB	2018-01-05	16580200.0
4	FB	2018-01-08	34574900.0

## Observations:

**OBV is cumulative**, meaning larger stocks like GOOG or AMZN may have extreme values due to high trading volume.

**Negative OBV values indicate downward momentum**, suggesting periods of sell pressure.

**Your OBV distribution aligns with high-volume stocks**, meaning your calculations are working as expected.

Moving Average Convergence Divergence (MACD)—one of the most popular momentum indicators used to identify trend direction and potential reversals.

Compute MACD Components

```
In [255... df_faang["EMA_12"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.ewm(
df_faang["EMA_26"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.ewm(
```

Compute MACD Line

```
In [256... df_faang["MACD"] = df_faang["EMA_12"] - df_faang["EMA_26"]
```

Compute Signal Line (9-day EMA of MACD)

```
In [258... df_faang["Signal_Line"] = df_faang.groupby("ticker")["MACD"].transform(lambda x: x.
```

Create DataFrame for analysis

```
In [259... macd_df = pd.DataFrame({
    "Ticker": df_faang["ticker"],
    "Date": df_faang.index, # Reference index as the date
    "MACD": df_faang["MACD"],
    "Signal Line": df_faang["Signal_Line"]
}).reset_index(drop=True)

# Display results
print(macd_df.describe()) # Summary statistics
print(macd_df.head()) # Preview first few rows
```

	Date	MACD	Signal Line
count	1255.000000	1255.000000	1255.000000
mean	627.000000	2.345486	2.587044
std	362.431603	18.391145	17.131601
min	0.000000	-82.597640	-67.632309
25%	313.500000	-4.249756	-3.982670
50%	627.000000	0.438979	0.404101
75%	940.500000	9.503180	8.589008
max	1254.000000	57.543031	49.816948

	Ticker	Date	MACD	Signal Line
0	FB	0	0.000000	0.000000
1	FB	1	0.259259	0.051852
2	FB	2	0.432306	0.127943
3	FB	3	0.763983	0.255151
4	FB	4	1.129211	0.429963



Possible improvement

```
In [260... df_faang["Date"] = pd.to_datetime(df_faang.index)
```

```
In [261... # Display results
print(macd_df.describe()) # Summary statistics
print(macd_df.head()) # Preview first few rows
```

	Date	MACD	Signal Line
count	1255.000000	1255.000000	1255.000000
mean	627.000000	2.345486	2.587044
std	362.431603	18.391145	17.131601
min	0.000000	-82.597640	-67.632309
25%	313.500000	-4.249756	-3.982670
50%	627.000000	0.438979	0.404101
75%	940.500000	9.503180	8.589008
max	1254.000000	57.543031	49.816948

Ticker	Date	MACD	Signal Line
0	FB	0 0.000000	0.000000
1	FB	1 0.259259	0.051852
2	FB	2 0.432306	0.127943
3	FB	3 0.763983	0.255151
4	FB	4 1.129211	0.429963

## Enough with question # 3

Lets move on to 4-9

### Exercie 4.

Crosstab for Earthquake Data

Load earthquake data (assuming 'tsunami', 'magType', and 'mag' columns exist)

```
In [274... earthquake_df = pd.read_csv("earthquakes.csv")
```

Display basic information to check available columns

```
In [275... print(earthquake_df.info())
print(earthquake_df.head()) # Preview first few rows
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9332 entries, 0 to 9331
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   mag             9331 non-null   float64
 1   magType         9331 non-null   object
 2   time            9332 non-null   int64
 3   place           9332 non-null   object
 4   tsunami         9332 non-null   int64
 5   parsed_place    9332 non-null   object
dtypes: float64(1), int64(2), object(3)
memory usage: 437.6+ KB
None
```

	mag	magType	time	place	tsunami	parsed_place
0	1.35	m1	1539475168010	9km NE of Aguanga, CA	0	California
1	1.29	m1	1539475129610	9km NE of Aguanga, CA	0	California
2	3.42	m1	1539475062610	8km NE of Aguanga, CA	0	California
3	0.44	m1	1539474978070	9km NE of Aguanga, CA	0	California
4	2.16	md	1539474716050	10km NW of Avenal, CA	0	California

Build crosstab showing max magnitude per tsunami/magType combo

```
In [277...] crosstab_df = earthquake_df.pivot_table(values="mag", index="tsunami", columns="magType")
```

Display the result

```
In [278...] print(crosstab_df)
```

magType	mb	mb_lg	md	mh	m1	ms_20	mw	mwb	mwr	mww
tsunami										
0	5.6	3.5	4.11	1.1	4.2	NaN	3.83	5.8	4.8	6.0
1	6.1	NaN	NaN	NaN	5.1	5.7	4.41	NaN	NaN	7.5

## Step 5:

Apply Rolling 60-Day Window

Apply 60-day rolling aggregations per ticker

```
In [284...] df_faang = pd.read_csv("faang.csv", parse_dates=["date"])
```

```
In [286...] # Apply 60-day rolling aggregations per ticker
df_faang["60D_Open"] = df_faang.groupby("ticker")["open"].transform(lambda x: x.rolling(60).max())
df_faang["60D_High"] = df_faang.groupby("ticker")["high"].transform(lambda x: x.rolling(60).max())
df_faang["60D_Low"] = df_faang.groupby("ticker")["low"].transform(lambda x: x.rolling(60).min())
df_faang["60D_Close"] = df_faang.groupby("ticker")["close"].transform(lambda x: x.rolling(60).max())
df_faang["60D_Volume"] = df_faang.groupby("ticker")["volume"].transform(lambda x: x.rolling(60).sum())

# Display results
print(df_faang[["ticker", "60D_Open", "60D_High", "60D_Low", "60D_Close", "60D_Volume"]])
```

	ticker	60D_Open	60D_High	60D_Low	60D_Close	60D_Volume
0	FB	NaN	NaN	NaN	NaN	NaN
1	FB	NaN	NaN	NaN	NaN	NaN
2	FB	NaN	NaN	NaN	NaN	NaN
3	FB	NaN	NaN	NaN	NaN	NaN
4	FB	NaN	NaN	NaN	NaN	NaN

```
In [287... print(df_faang[["ticker", "60D_Close"]].dropna().head(10))
```

	ticker	60D_Close
59	FB	179.880499
60	FB	179.519999
61	FB	179.031999
62	FB	178.561666
63	FB	178.032499
64	FB	177.550166
65	FB	177.038999
66	FB	176.540499
67	FB	176.161666
68	FB	175.944166

## Step 6 Building the Pivot Table

Create pivot table with FAANG stock averages

```
In [288... pivot_df = df_faang.pivot_table(
    index="ticker",
    values=["open", "high", "low", "close", "volume"],
    aggfunc="mean"
)
```

Display results

```
In [289... print(pivot_df)
```

	close	high	low	open	volume
ticker					
AAPL	47.263357	47.748526	46.795877	47.277859	1.360803e+08
AMZN	1641.726176	1662.839839	1619.840519	1644.072709	5.648994e+06
FB	171.510956	173.613347	169.303148	171.472948	2.765860e+07
GOOG	1113.225134	1125.777606	1101.001658	1113.554101	1.741965e+06
NFLX	319.290319	325.219322	313.187330	319.620558	1.146962e+07

## Next up is Step 7:

Calculating Z-Scores for Amazon's Data in Q4 2018

```
In [290... import pandas as pd
from scipy.stats import zscore
```

Filter FAANG data for Amazon (AMZN) in Q4 2018

```
In [291... amzn_q4 = df_faang[(df_faang["ticker"] == "AMZN") & (df_faang["date"] >= "2018-10-01")]
```

Display the first few rows

In [292... `print(amzn_q4.head())`

	ticker	date	high	low	open	close \
690	AMZN	2018-10-01	2033.189941	2003.599976	2021.989990	2004.359985
691	AMZN	2018-10-02	2013.390015	1965.770020	1999.989990	1971.310059
692	AMZN	2018-10-03	1989.699951	1949.810059	1981.699951	1952.760010
693	AMZN	2018-10-04	1956.000000	1896.569946	1949.000000	1909.420044
694	AMZN	2018-10-05	1929.079956	1862.829956	1917.989990	1889.650024

	volume	60D_Open	60D_High	60D_Low	60D_Close	60D_Volume
690	3460500.0	1894.282505	2050.5	1716.229980	1894.499168	281006500.0
691	5400700.0	1898.881504	2050.5	1731.000000	1898.370669	283395200.0
692	5253100.0	1902.934336	2050.5	1734.000000	1901.865503	285645400.0
693	7257000.0	1906.451170	2050.5	1739.319946	1904.439170	289692600.0
694	6822300.0	1909.009170	2050.5	1739.319946	1905.989671	291982200.0

**Apply Z-score normalization to numeric columns (excluding ticker and date)**

In [293... `amzn_zscore = amzn_q4.drop(columns=["ticker", "date"]).apply(zscore)`

Display results

In [295... `print(amzn_zscore.head())`

	high	low	open	close	volume	60D_Open	60D_High \
690	2.387026	2.522211	2.356591	2.405011	-1.643506	0.842151	0.523572
691	2.245193	2.265485	2.208392	2.172347	-0.868802	0.900937	0.523572
692	2.075493	2.157176	2.085185	2.041758	-0.927738	0.952743	0.523572
693	1.834088	1.795871	1.864908	1.736654	-0.127599	0.997696	0.523572
694	1.641251	1.566901	1.656015	1.597477	-0.301170	1.030394	0.523572

	60D_Low	60D_Close	60D_Volume
690	1.506719	0.908550	-1.733558
691	1.616224	0.957078	-1.686799
692	1.638466	1.000885	-1.642751
693	1.677908	1.033146	-1.563528
694	1.677908	1.052581	-1.518709

## Observations & Insights

Price metrics (high, low, open, close) show consistently high Z-scores (above 1.5 in early October), indicating Amazon's prices were significantly above its mean value during this period.

Trading volume has negative Z-scores, meaning Amazon's trading activity was lower than its typical average, suggesting less market participation during these days.

60-day rolling prices show relatively stable Z-scores (~0.9 to 1.05), reinforcing the idea that Amazon's performance over longer periods was fairly consistent.

Volume Z-score trends upwards (less negative), meaning Amazon's trading activity picked up later in the period.

## Step 8a:

Adding Event Descriptions. We'll create a DataFrame with ticker, date, and event details for Facebook (FB), then set the proper index for merging.

Create event DataFrame

```
In [298... events_df = pd.DataFrame({
    "ticker": ["FB", "FB", "FB"],
    "date": ["2018-07-25", "2018-03-19", "2018-03-20"],
    "event": ["Disappointing user growth announced after close.",
              "Cambridge Analytica story",
              "FTC investigation"]
})
```

Convert date column to datetime format

```
In [299... events_df["date"] = pd.to_datetime(events_df["date"])
```

Display DataFrame

```
In [300... print(events_df)
```

	ticker	date	event
0	FB	2018-07-25	Disappointing user growth announced after close.
1	FB	2018-03-19	Cambridge Analytica story
2	FB	2018-03-20	FTC investigation

Setting the Index for Merging

Set the index to ['date', 'ticker']

```
In [301... events_df.set_index(["date", "ticker"], inplace=True)
```

Display the updated DataFrame

```
In [302... print(events_df)
```

	date	ticker	event
	2018-07-25	FB	Disappointing user growth announced after close.
	2018-03-19	FB	Cambridge Analytica story
	2018-03-20	FB	FTC investigation

Merge Event Data with FAANG Data

Merge FAANG data with event descriptions using an outer join

```
In [303... df_faang_events = df_faang.merge(events_df, how="outer", left_on=["date", "ticker"])
```

Display merged dataset

```
In [304... print(df_faang_events.head())
```

	ticker	date	high	low	open	close	\
251	AAPL	2018-01-02	43.075001	42.314999	42.540001	43.064999	
502	AMZN	2018-01-02	1190.000000	1170.510010	1172.000000	1189.010010	
0	FB	2018-01-02	181.580002	177.550003	177.679993	181.419998	
1004	GOOG	2018-01-02	1066.939941	1045.229980	1048.339966	1065.000000	
753	NFLX	2018-01-02	201.649994	195.419998	196.100006	201.070007	

	volume	60D_Open	60D_High	60D_Low	60D_Close	60D_Volume	event
251	102223600.0	NaN	NaN	NaN	NaN	NaN	NaN
502	2694500.0	NaN	NaN	NaN	NaN	NaN	NaN
0	18151900.0	NaN	NaN	NaN	NaN	NaN	NaN
1004	1237600.0	NaN	NaN	NaN	NaN	NaN	NaN
753	10966900.0	NaN	NaN	NaN	NaN	NaN	NaN

Verify Events Are Merged

Show rows where Facebook events should be present

```
In [306... print(df_faang_events[df_faang_events["event"].notna()])
```

	ticker	date	high	low	open	close	\
52	FB	2018-03-19	177.169998	170.059998	177.009995	172.559998	
53	FB	2018-03-20	170.199997	161.949997	167.470001	168.149994	
141	FB	2018-07-25	218.619995	214.270004	215.720001	217.500000	

	volume	60D_Open	60D_High	60D_Low	60D_Close	\
52	88140100.0	NaN	NaN	NaN	NaN	
53	129851800.0	NaN	NaN	NaN	NaN	
141	58954200.0	193.029334	218.619995	170.229996	193.826667	

	60D_Volume	event
52	NaN	Cambridge Analytica story
53	NaN	FTC investigation
141	1.099926e+09	Disappointing user growth announced after close.

## Step 9:

Index Transformation with transform(). This method will represent all stock values in terms of the first available date, helping visualize growth over time.

Normalize all values relative to the first date per ticker

Normalize all values relative to the first date per ticker (excluding 'date')

```
In [309... df_faang_indexed = df_faang.set_index("date").groupby("ticker").transform(lambda x:
```

Display indexed dataset

```
In [310... print(df_faang_indexed.head())
```

	high	low	open	close	volume	60D_Open \
date						
2018-01-02	1.000000	1.000000	1.000000	1.000000	1.000000	NaN
2018-01-03	1.017623	1.021290	1.023638	1.017914	0.930294	NaN
2018-01-04	1.025498	1.036891	1.040635	1.016040	0.764708	NaN
2018-01-05	1.029298	1.041566	1.044518	1.029931	0.747828	NaN
2018-01-08	1.040313	1.049451	1.053579	1.037813	0.991340	NaN


	60D_High	60D_Low	60D_Close	60D_Volume
date				
2018-01-02	NaN	NaN	NaN	NaN
2018-01-03	NaN	NaN	NaN	NaN
2018-01-04	NaN	NaN	NaN	NaN
2018-01-05	NaN	NaN	NaN	NaN
2018-01-08	NaN	NaN	NaN	NaN

## Plot Indexed Price Trends

### Select relevant columns for visualization

```
In [311... df_faang_indexed[["high", "low", "open", "close"]].plot(figsize=(12, 6), title="FAA
```

```
Out[311... <Axes: title={'center': 'FAANG Indexed Price Trends'}, xlabel='date'>
```



Key Takeaways from Your FAANG Analysis Earthquake Crosstab: Mapped tsunami occurrences to magnitude types, revealing trends in seismic activity.

FAANG Rolling 60-Day Aggregations: Smoothed stock data to uncover long-term patterns in Open, High, Low, Close, and Volume.

Pivot Table for FAANG Comparison: Compared major tech stocks, highlighting price levels and trading volumes.

Amazon's Q4 2018 Z-Scores: Standardized values to track deviations and detect performance shifts.

Event-Driven Stock Analysis: Merged Facebook's major events (Cambridge Analytica, FTC probe, user growth slowdown) with its stock data to analyze investor reactions.

Index Transformation: Converted stock prices into relative terms, revealing overall growth trends.

Visualization: Brought the insights to life with plots, making patterns clearer and more intuitive.